# Lusail: A System for Querying Linked Data at Scale

Ibrahim Abdelaziz* Essam Mansour◇ Mourad Ouzzani◇ Ashraf Aboulnaga◇ Panos Kalnis*
*King Abdullah University of Science & Technology ◇Qatar Computing Research Institute, HBKU
{first}.{last}@kaust.edu.sa, {emansour,mouzzani,aaboulnaga}@hbku.edu.qa

## ABSTRACT

The RDF data model allows different organizations to publish datasets that are interlinked to other RDF datasets, where each dataset is independently maintained and queryable via a SPARQL endpoint. This creates a large decentralized geo-distributed queryable RDF graph, and many applications would benefit from querying this graph through a federated SPARQL query processor. A crucial factor for good performance in federated query processing is pushing as much computation as possible to the local endpoints. Surprisingly, existing federated SPARQL engines are not effective at this task, since they rely only on schema information to push computation to the endpoints, and schema information is limited in its power. Therefore, these systems often cause unnecessary data retrieval and communication, leading to poor scalability and response time. This paper addresses these limitations and presents *Lusail*, a scalable and efficient federated SPARQL system for querying large geo-distributed RDF graphs through different endpoints. Lusail uses a novel query rewriting algorithm to push computation to the local endpoints by relying on information about the RDF instances and not only the schema. The query rewriting algorithm has the additional advantage of exposing parallelism in query processing, which Lusail exploits through advanced scheduling at query run time. Our experiments on billions of triples of real and synthetic data show that Lusail outperforms state-of-the-art systems by orders of magnitude in terms of scalability and response time.

## 1. INTRODUCTION

The Resource Description Framework (RDF) is extensively used to represent structured data on the Web. RDF uses a simple graph data model in which the data are in the form of triples ⟨`subject`, `predicate`, `object`⟩. A key feature is the ability to link two entities from two different RDF datasets which are maintained by two independent authorities, as shown in Figure 1. Through such links, *large decentralized graphs* are created among a large number

Figure 1: Decentralized graphs for two universities managed by independent geo-distributed SPARQL endpoints (EP). The red dotted line represents an interlink between endpoints, i.e., a vertex in an endpoint refers to another vertex in another endpoint. Thus, to get the address of the university from which Tim got his PhD, the interlink from EP2 to EP1 has to be traversed.

```
SELECT ?S ?P ?U ?A WHERE{
 ?S ub:advisor ?P .        ?S rdf:type ub:graduateStudent .
 ?P ub:teacherOf ?C .      ?P rdf:type ub:associateProfessor .
 ?S ub:takesCourse ?C .    ?C rdf:type ub:graduateCourse .
 ?P ub:PhDDegreeFrom ?U . ?U ub:address ?A . }
```

Figure 2: $Q_a$: A SPARQL query over decentralized RDF graphs across different universities. This query has to traverse the interlink between EP2 and EP1.

of geo-distributed RDF stores where each RDF store can be queried through its own SPARQL endpoint. By a recent count, decentralized RDF graphs consist of more than 85 billions triples in more than 3400 datasets[1] in different domains, such as media, government, and life sciences[2] [27].

Users can retrieve data from an individual dataset by issuing SPARQL queries against the SPARQL endpoint of this dataset. However, it is often very useful to issue SPARQL queries where each query integrates data from multiple RDF datasets, which requires *federated query processing* of SPARQL. For example, Figure 2 shows a query ($Q_a$) on data from the LUBM benchmark [12] at two endpoints. $Q_a$ returns all students who are taking courses with their advisors along with the URI and location of the advisors' alma mater. $Q_a$ has three answers: (Kim, Joy, CMU, "CCCC"), (Kim, Tim, MIT, "XXX") and (Lee, Ben, MIT, "XXX"). One cannot simply evaluate $Q_a$ independently at each endpoint

---
[1]http://stats.lod2.eu/
[2]http://bio2rdf.org/

and concatenate their results as this will miss the results about Tim since EP2 does not have the address of MIT. Instead, we need a federated query processor that can automatically identify the endpoints that can answer each triple pattern, and that can detect interlinks between endpoints and automatically traverse them as needed to compute the query answer. A federated query processor would decompose $Q_a$ into subqueries, send each subquery to the relevant endpoint, and compute the answer to $Q_a$ from the results of the subqueries. Computing the answer to a query typically requires the federated query processor to *join* the results obtained from the endpoints. For example, in $Q_a$ this join would combine the address of MIT from EP1 with the information about Tim from EP2.

Conceptually, a query like $Q_a$ can be processed by sending each of its triple patterns to all the endpoints, retrieving all matching triples from the endpoints, and joining all of these triples at the federated query processor to compute the query answer. This strategy is clearly inefficient since it sends triple patterns to endpoints even if they have no answers for them, retrieves triples that may not be relevant (e.g., it would retrieve the address of MIT in $Q_a$ even if no advisors have MIT as their alma mater), and joins triples at the federated query processor even if they could be joined at the endpoints. Thus, this strategy would result in an unnecessarily large number of requests to the endpoints and unnecessarily large amounts of data retrieved from the endpoints and transferred over the network to the federated query processor. To avoid these unnecessary overheads, it is important for federated query processors to push as much processing as possible to the local SPARQL endpoints.

Existing SPARQL federated query processing systems rely on schema information to push processing to the endpoints. For example, they use SPARQL ASK queries to check whether or not a triple pattern has an answer at an endpoint. If a group of triple patterns can be answered exclusively by one endpoint, then it is possible to send this group to the endpoint as one unit, known as an *exclusive group* [29]. Relying solely on schema information is not effective since RDF sources often utilize similar ontologies (e.g., EP1 and EP2 in Figure 2 have the same predicates), thus a triple pattern could be answerable by multiple endpoints and therefore cannot be part of an exclusive group. In this case, the triple pattern is sent to all the endpoints that can answer it and the values in the retrieved triples are bound to other triple patterns, and these triple patterns with bound values are sent to the endpoints to retrieve further triples. This is known as a *bound join* operation, and effectively amounts to the query being processed one triple pattern at a time. This strategy retrieves unnecessary data from the endpoints, since it retrieves all data matching a triple pattern even if this data is not useful for the rest of the query. Moreover, this process limits the available parallelism since only one join step can be processed at a time, and the federated query processor has to wait for the results of this join step before issuing the next join.

This paper addresses the limited ability of existing systems to push query processing to the local endpoints. We present Lusail, a system for federated SPARQL query processing over decentralized RDF graphs that achieves scalability and efficiency through better query decomposition. A key defining feature of Lusail is that it is the first system to decompose the federated query based on *instance* information not just schema information. That is, Lusail decomposes the query based on knowledge of the locations of the actual RDF triples matching triple patterns in the query. This knowledge helps us identify, for example, that the instances matching the variable ?S in ⟨?S, ub:advisor, ?P⟩ and ⟨?S, ub:takesCourse, ?C⟩ in $Q_a$ are always located in the same endpoint, so these triple patterns can be joined locally at the endpoint even though schema information tells us that both endpoints can answer both triple patterns. In contrast, the instances matching the variable ?U in ⟨?P, ub:PhDDegreeFrom, ?U⟩ and ⟨?U, ub:address, ?A⟩ are sometimes located in different endpoints, so these triple patterns cannot be joined locally.

Lusail processes queries in a two-phase strategy: (*i*) Locality-Aware DEcomposition (LADE) of the query into subqueries to maximize the computations at the endpoints and minimize intermediary results, and (*ii*) Selectivity-Aware and Parallel Execution (SAPE) to reduce network latency and increase parallelism. When two or more triple patterns have solutions at different endpoints, LADE investigates at each endpoint which triple patterns access remote data based on the actual location of the instances satisfying these patterns. Subsequently, LADE decomposes the query into a set of subqueries that will be executed independently at one or more endpoints. Afterward, SAPE dynamically decides to delay subqueries expected to return large results. It also chooses the join order among the subquery results based on their actual size and the highest degree of parallelism that could be achieved. SAPE uses a cost model for balancing between remote requests and local computations.

We have demonstrated Lusail in [18] and discussed the challenges of processing federated SPARQL queries at scale on a poster in [1]. In this paper, we describe the full system. In summary, our contributions are as follows:

- We study different federated SPARQL query processing approaches in terms of their scalability limitations w.r.t data sizes and number of endpoints. (Section 2)

- A locality-aware decomposition method that dramatically reduces the number of remote requests and allows for better utilization of the endpoints. We also provide a proof of correctness. (Section 4)

- A cost model that uses lightweight runtime statistics to decide the order of submitting subqueries and the tree execution plan for joining the results of these subqueries in a non-blocking fashion. This leads to a parallel execution that balances between remote requests and local computations. (Section 5)

- Our experiments on real data and synthetic benchmarks with billions of triples show that Lusail outperforms state-of-the-art systems by up to three orders of magnitude and scales-up to 256 endpoints compared to 4 endpoints in existing systems. (Section 6)

We present the architecture of Lusail in Section 3, discuss related work in Section 7, and conclude in Section 8.

## 2. LIMITATIONS TO SCALABILITY

In this section, we examine the performance of FedX [29], a federated query processor that was shown to outperform other similar systems [25], as the number of endpoints increases. FedX starts processing a query by sending ASK
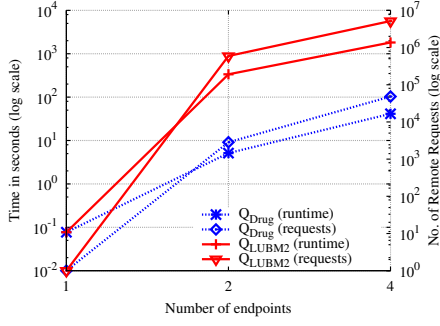
**Figure 3: Sensitivity of FedX to varying endpoints.**

requests to all the endpoints to determine the relevant endpoints. The result of this source selection step is cached for futre invocations of the query. FedX then optimizes the join order and identifies exclusive groups. It uses block nested loops join to reduce the number of HTTP requests in the bound join. We use the QFed [23] and LUBM [12] benchmarks with two queries, the *Drug* query from QFed and *Q2* from LUBM. *Drug* finds all medicines that target asthma and obtains information about them. It uses up to 4 datasets (in 4 endpoints). *Q2* finds graduate students who are registered in courses delivered by their advisor. Each endpoint in the LUBM benchmark corresponds to a university. We run each query multiple times so that FedX can cache the results of source selection.

Figure 2 shows the response time and the number of remote HTTP requests for different numbers of endpoints, excluding the ASK queries made during source selection. We see that both response time and number of requests increase by orders of magnitude as the number of endpoints increases. Processing one triple pattern at a time while binding query variables to values from intermediate results causes a huge number of remote requests, which limits scalability. Even with the optimizations employed by FedX, remote requests are still a bottleneck. In the bound join, the number of remote requests depends on the size of the intermediate data. The bound join also reduces parallelism since only one join step is processed at a time. We thus need a way to avoid performing the bound join one triple pattern at a time.

## 3. THE LUSAIL ARCHITECTURE

The Lusail architecture is shown in Figure 4. Lusail evaluates the input queries using a set of geo-distributed SPARQL endpoints. It analyzes each query to identify the relevant endpoints and its correct decomposition that achieves high parallelism and minimal communication cost. After that, Lusail sends the subqueries to the relevant endpoints, joins their results, and sends the query answer back to the user.

**Locality-Aware Decomposition (LADE):** Query decomposition starts by analyzing the query to identify the relevant endpoints (source selection). Like similar systems, we use a set of SPARQL ASK queries, one for each triple pattern, and we cache the results. However, in contrast to other systems, LADE takes the additional step of checking, for each pair of triple patterns with a common (or join) variable, whether the pair can be evaluated as one unit by the relevant endpoints. To do so, LADE utilizes the knowledge
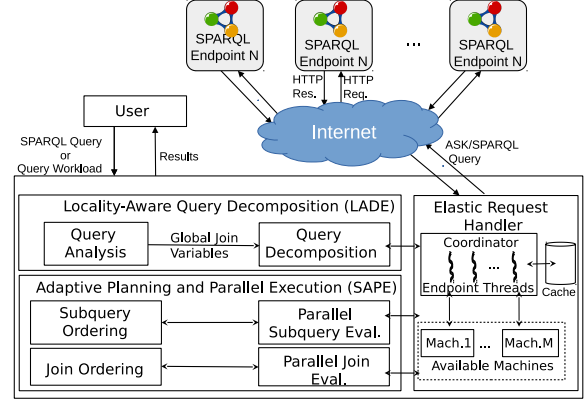


**Figure 4: The Lusail system architecture.**

of the locations of the actual RDF triple instances matching a query variable. The result of this check determines a group of triple patterns, i.e., a subquery, that can be sent together to an endpoint. Based on this analysis, LADE decomposes the query into a set of independent subqueries.

**Selectivity-Aware Planning and Parallel Execution (SAPE):** SAPE takes as input the set of subqueries produced by LADE and schedules them for execution. This set of independent subqueries can be submitted concurrently for execution at each of the relevant endpoints, and Lusail can use one thread per endpoint to collect their results. SAPE uses cardinality estimates for the different triple patterns to delay subqueries that are expected to return large results. The results of these subqueries will then need to be joined by SAPE using a parallel join, where the join order is determined based on the actual sizes of the subquery results. SAPE achieves a high degree of parallelism while minimizing the communication cost by ($a$) obtaining results from different endpoints simultaneously, and ($b$) utilizing different threads in joining the results.

**Elastic Request Handler (ERH):** Lusail utilizes multiple threads for evaluating the ASK queries from LADE or the subqueries from SAPE at the endpoints. ERH manages the allocation of threads from one or more machines to these tasks, where the number of available threads is determined by the number of physical cores.

## 4. LOCALITY-AWARE DECOMPOSITION

To push as much processing to the endpoints as possible, LADE maximizes the number of triple patterns in a given query that can be sent together to each endpoint. In decentralized RDF graphs, data instances matching a pair of triples may not be located in the same endpoint, e.g., the triples having `?U` as a common variable in Figure 1. Thus, putting this pair in the same subquery may miss results. LADE starts by analyzing which triples cannot be in the same subquery, and identifying the common variables in these triples as *global join variables (GJV)*. Then, it decomposes the conjunction of triple patterns into subqueries. We assume no prior knowledge of the data sources, such as schema, data distribution, or statistics. LADE relies solely on a set of check queries written in SPARQL.
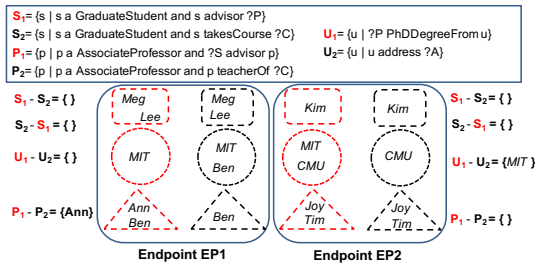
$S_1$= {s | s a GraduateStudent and s advisor ?P}
$S_2$= {s | s a GraduateStudent and s takesCourse ?C}    $U_1$= {u | ?P PhDDegreeFrom u}
$P_1$= {p | p a AssociateProfessor and ?S advisor p}    $U_2$= {u | u address ?A}
$P_2$= {p | p a AssociateProfessor and p teacherOf ?C}

$S_1$ - $S_2$= { }     Meg     Meg     $S_1$ - $S_2$= { }
                       Lee     Lee     Kim     Kim
$S_2$ - $S_1$ = { }            MIT             $S_2$ - $S_1$ = { }
                               Ben     MIT     CMU
$U_1$ - $U_2$ = { }    MIT            CMU       Joy     $U_1$ - $U_2$ = {MIT}
                       Ben            Tim
                       Ann     Ben     Joy     Joy
$P_1$ - $P_2$ = {Ann}  Ben            Tim     Tim      $P_1$ - $P_2$ = { }

**Endpoint EP1**          **Endpoint EP2**

**Figure 5: Locality analysis of data instances in EP1 and EP2 from Figure 1 that match ?S, ?U, and ?P in a pair of triples in $Q_a$.**

```
1 SELECT ?P WHERE {
2 ?P rdf:type T
3 ?S < Predicate_i > ?P .
4 FILTER NOT EXISTS { SELECT ?P WHERE {
5   ?P < Predicate_j > ?C .
6 }} . } LIMIT 1
```

**Figure 6: A Lusail SPARQL check query to detect whether ?P is a global join variable or not. The check query returns zero or only one value.**

## 4.1 Detecting Global Join Variables

A global join variable ($v$) is a variable that appears in at least two different triple patterns such that these triple patterns, when taken together, cannot be solved by a single endpoint. A global join between data coming from two or more endpoints will be needed. Given two triple patterns, $TP_i$ and $TP_j$, in a subquery, a GJV may appear in the triple patterns as: (*i*) *object* in $TP_i$ and *subject* in $TP_j$, (*ii*) *object* in both patterns, or (*iii*) *subject* in both. Let $v_i$ and $v_j$ be the sets of instances of $v$ that satisfy $TP_i$ and $TP_j$, respectively.

$Q_a$ (Figure 2) has four variables appearing in more than one triple pattern, namely ?S, ?U, ?P, and ?C. Figure 5 shows our analysis for the first three variables. In EP1 and EP2, all instances matching ?S in ⟨?S, ub:advisor, ?P⟩ are co-located with all instances matching ?S in ⟨?S, ub:takesCourse, ?C⟩. Thus, ?S is not a GJV and hence the two corresponding triple patterns can be sent together in a single subquery to each relevant endpoint. However, for the triples involving ?U, ⟨?P, ub:PhDDegreeFrom, ?U⟩ and ⟨?U, ub:address, ?A⟩, we notice that in EP2 there is a professor, Tim, who got his PhD from another university. Thus, to get the address of that university, we need to perform a join between data fetched from EP1 and EP2. Therefore, ?U is a GJV.

We now describe how LADE detects GJVs by determining the actual location of data instances depending on the roles they play, i.e., object or subject. We first discuss how to merge two triple patterns and then generalize to more than two (see Algorithm 1). Two triple patterns $TP_i$ and $TP_j$ are put together in a single subquery under two conditions: (*i*) both triple patterns have the same list of relevant endpoints, and (*ii*) each relevant endpoint can fully answer both triple patterns without missing any result, i.e, all instances that match $v$ in $TP_i$ and $TP_j$ are in the same endpoint.
**Object and Subject.** Consider the variable ?U in $Q_a$ (Figure 2). It appears as an object in $TP_i$: ?P ub:PhDDegreeFrom ?U and as a subject in $TP_j$: ?U ub:address ?A. Checking the location of the data instances $v_i$ and $v_j$ that match ?U in each endpoint has two cases: (*i*) remote instances, where $v_i$ and $v_j$ are located in different endpoints, i.e., all or some professors received their PhD from another university (in a different endpoint); e.g., EP2 in Figure 5, and (*ii*) local instances, where all $v_i$ and $v_j$ are located in the same endpoint, i.e., all professors teaching in a university $A$ received their PhD from $A$ (in the same endpoint), e.g., EP1 in Figure 5.

We check the relative complement (i.e., set difference) of $v_i$ and $v_j$ in all relevant endpoints by sending a SPARQL query to each endpoint. If at least one of these endpoints has instances in $v_i$ but not in $v_j$, then $v$ is a GJV. At each endpoint, we check for each data instance appearing as an object in $TP_i$ whether this instance appears locally as a subject in $TP_j$. Once a common variable is found to be a GJV, the triple patterns cannot be combined in the same subquery even for those endpoints that return an empty result for the difference in the instances, e.g., the pair of triples where ?U is a common variable in Figure 5. This allows us to have simple plans and better parallel execution.

Set difference ($-$) is implemented using *FILTER NOT EXISTS* (Figure 6) where $TP_i$: ⟨?S, Predicate$_i$, ?P⟩, and $TP_j$: ⟨?P, Predicate$_j$, ?C⟩. If there is a triple pattern setting a type for $v$ (⟨?P, rdf:type, $T$⟩), we use it to limit the check to only the relevant values of $v$. Since Lusail needs to only know whether the result is an empty set, we use *LIMIT 1*. Note that our goal is to check whether there are data instances matching the predicate $j$ regardless of the actual values. So, if the query specifies a constant in the triple pattern mentioned in the filter of the check query (i.e., 'XXX'), Lusail will replace it with a variable in the check query.
**Objects/Subjects Only.** If a variable appears only as *object*, respectively *subject*, in both triple patterns $TP_i$ and $TP_j$, Lusail checks in each relevant endpoint that $v_i - v_j$ and $v_j - v_i$ are both empty. As shown in Figure 5, the variable ?S appears as subject in both ⟨?S, ub:advisor, ?P⟩ and ⟨?S, ub:takesCourse, ?C⟩. Having two empty sets in the same endpoint means that (*i*) any graduate student ?S having an advisor ?P should take a course ?C and (*ii*) any graduate student ?S taking a course ?C should have an advisor ?P, all located in the same endpoint.

Algorithm 1 receives a query and a list of relevant endpoints and outputs a set of GJVs ($V$) along with the triple patterns that caused each variable to be a GJV. It assumes that source selection is already done using ASK requests or the Lusail cache. The algorithm starts by retrieving the set of query variables and triple patterns. Each variable is associated with its subject and object patterns (line 2). The algorithm iterates over the variables to detect GJVs.

If the variable joins triple patterns from different sources, then it is a GJV (lines 8-11). There is no need to check the other conditions. Otherwise, Algorithm 1 formulates a set of check queries as discussed above. For the *object only* or *subject only* cases, Algorithm 1 formulates check queries for all possible pairwise combinations of the triple patterns associated with the variable (lines 13-14). For the *object and subject* case, the check query is a combination of object triples and subject triples (lines 15-16).

The algorithm uses the elastic request handler (Figure 4) to execute check queries. It initializes the handler with the size of the thread pool and the set of endpoints (line 18). Then, it iterates over all check queries and executes each

**Algorithm 1:** Detecting Global Join Variables

**Input**: Input query ($Q$), Set of relevant sources ($Sources$)
**Result**: List of Global Join Variables ($V$)

1   $Triples \leftarrow Q.\texttt{getTriplePatterns}();$
2   $vars \leftarrow \texttt{getJoinEntities}\ (Triples);$
3   $chkQueries \leftarrow \emptyset;$
4   $V \leftarrow \emptyset;$
5   **foreach** $var_i$ *in* $vars$ **do**
6      $pairWiseTriples \leftarrow \texttt{getPairTriples}\ (var_i.Triples);$
7      $joinVar \leftarrow$ False;
8      **foreach** $pair_i$ *in* $pairWiseTriples$ **do**
9        **if** $pair_i[0].sources \neq pair_i[1].sources$ **then**
10          $V.\texttt{addJoinVar}\ (var_i.varName,\ pair_i);$
11          $joinVar \leftarrow$ True;
12      **if** $joinVar$ *is True* **then** continue ;
13      **if** $var_i$ *is subject only* $||$ $var_i$ *is object only* **then**
14        $chkQueries \leftarrow \texttt{formulatePairWiseQuery}\ (var_i,$
       $pairWiseTriples\ );$
15      **if** $var_i$ *is subject and object* **then**
16        $chkQueries \leftarrow \texttt{formulateSubjObjQuery}\ (var_i,$
       $var_i.\texttt{subjTriples},\ var_i.\texttt{objTriples});$
17   **if** $chkQueries$ *is not empty* **then**
18      $ReqHandler \leftarrow \texttt{initializeRequestHandler}\ (thrdPoolSize,$
     $Sources);$
19      **foreach** $chkQry_i$ *in* $chkQueries$ **do**
20        //each $chkQry_i$ is attached with its relevant sources;
21        $RES = ReqHandler.\texttt{executechkQAtRelSrcs}\ (chkQry_i);$
22        **if** $RES$ *is not empty* **then**
23          $V.\texttt{addJoinVar}\ (chkQry_i.varName,$
         $chkQry_i.triples);$
24   **return** $V;$

---

**Algorithm 2:** Query Decomposition

**Input**: Input query ($Q$), set of GJVs ($V$)
**Result**: Set of independent *subqueries*

1   $bestDecomposition \leftarrow \emptyset;$
2   $minDecompCost \leftarrow infinity;$
3   **if** $V$ *is empty* **then** $subqueries.\texttt{add}\ (Q);$
4   $Triples \leftarrow Q.\texttt{getTriplePatterns}();$
5   **foreach** $jvar$ *in* $V$ **do**
6      $visitedTriples \leftarrow \emptyset;$
7      $nodes \leftarrow \emptyset;$
8      $subqueries \leftarrow \emptyset;$
9      $nodes.\texttt{push}\ (jvar_i);$
10     **while** $nodes$ *is not empty* **do**
11        $vrtx \leftarrow nodes.\texttt{pop}\ ();$
12        $edges \leftarrow vrtx.\texttt{edges}\ ();$
13        **if** $subqueries$ *is empty* **then**
14          **foreach** $edge_i$ *in* $edges$ **do**
15            **if** $\texttt{visited}\ (edge_i,\ visitedTriples)$ **then**
           continue ;
16            $sq \leftarrow \texttt{createSubquery}(edge_i);$
17            $subqueries.\texttt{add}\ (sq);$
18            $nodes.\texttt{push}\ (edge_i.destNode);$
19            $visitedTriples.\texttt{add}\ (edge_i);$
20          continue;
21        $parentSq \leftarrow \texttt{getParentSubquery}\ (vrtx, subqueries);$
22        **foreach** $edge_i$ *in* $edges$ **do**
23          **if** $\texttt{visited}\ (edge_i,\ visitedTriples)$ **then** continue ;
24          **if** $\texttt{canBeAddedToSubQ}\ (parentSq,\ edge_i,\ V)$ **then**
25            $parentSq \leftarrow \texttt{addToSubquery}\ (parentSq,\ edge_i);$
26          **else**
27            $sq \leftarrow \texttt{createSubquery}(edge_i);$
28            $subqueries.\texttt{add}\ (sq);$
29          $nodes.\texttt{push}\ (edge_i.destNode);$
30          $visitedTriples.\texttt{add}\ (edge_i);$
31      **if** $visitedTriples \equiv Triples$ **then**
32        $subqueries \leftarrow \texttt{mergeSubQ}\ (subqueries);$
33        $cost \leftarrow \texttt{estimateCost}\ (subqueries);$
34        **if** $cost < minDecompCost$ **then**
35          $bestDecomposition \leftarrow subqueries;$
36          $minDecompCost \leftarrow cost;$
37   **return** $bestDecomposition;$

---

at the relevant endpoints (lines 19-23). If the query returns any results, then the corresponding variable is a GJV (lines 22-23). The algorithm returns the set of GJVs along with the triple patterns that caused each variable to be a GJV.

Assuming that $|V|$ is the number of variables appearing in more than one triple pattern in the query and $|T|$ is the number of triples, the maximum number of check queries, $C_Q$, is bound by $O(|V| * |T|^2)$. Assuming $N$ relevant endpoints, LADE creates a maximum of $N * C_Q$ HTTP requests. Since the number of triple patterns in real-world SPARQL queries is usually small [10], the number of GJVs is also small. Therefore, $N * C_Q$ will be typically small. In addition, the check queries are lightweight and have a minimal overhead (see Section 6.4).

## 4.2   Query Decomposition

Algorithm 2 decomposes a query $Q$ into multiple subqueries to be sent to different endpoints. If $Q$ has no GJVs, the algorithm returns $Q$ (line 3). Otherwise, LADE uses the set of GJVs and the source selection information to decompose $Q$. It iterates over all join variables in any order using the current join variable as a root. It tries to find the best decomposition that leads to a set of subqueries with minimal execution cost (cost estimation is discussed in Section 5). The algorithm has two phases: branching (lines 9-30) and merging (lines 32).

In the branching phase, we build a query tree with the current join variable as its root (line 9). An initial set of subqueries is created at the root, one subquery per child (lines 13-20) and each subquery is expanded through depth first traversal (lines 21-30). A triple pattern is added (lines 23-24) if both the subquery and the triple pattern have the same relevant sources, and the addition of the pattern does not cause a query variable to be a GJV. If one of the condi-

tions is invalid, a new subquery is created from the current triple pattern and added to the set of subqueries (lines 27-28). In both cases, the edge destination node is added to the nodes stack and marked as visited (lines 29-30).

The merging phase (line 32) starts once all triple patterns are assigned to one of the subqueries (line 31). The function $mergeSubQ$ (line 32) loops through the set of subqueries and merges a pair of subqueries if they have common variables, the same relevant sources, and no pair of triple patterns from both subqueries has a common variable that is global. If the estimated cost (line 33) of the current decomposition is less than other decompositions, the algorithm updates the minimum cost value and selects the current decomposition as the current best. The algorithm continues to check the other possible decompositions using the remaining join variables. The algorithm return the best subqueries decomposition (line 37). For simplicity, the pseudo-code of the algorithm assumes a connected query graph. It can be easily adjusted for the general case.

Figure 7 shows two possible decompositions for $Q_a$ (Figure 2), which has two GJVs, namely ?U and ?P. The generated set of subqueries may change depending on the order in which variables are selected during query decomposition (line 5 in Algorithm 2). However, all decompositions produce the same result set and do not miss any triple (Section 4.4 for detail), but some decompositions may generate

**Figure 7: Two possible decompositions of $Q_a$, where the GJVs are ?U and ?P. Any pair of predicates, which causes a variable to be a GJV, cannot be in the same subquery.**

more intermediate results and thus cost more. To avoid a costly decomposition, LADE enumerates all possible decompositions and chooses at compile time the best decomposition expected to minimize the intermediate results.

Decomposing the query into subqueries is similar to a depth first traversal. Its runtime complexity is $O(|V| + |T|)$ where $|V|$ is the total number of query variables and $|T|$ is the number of triple patterns. The outer loop (line 5 in Algorithm 2) can repeat the traversal for few more times if there are triple patterns that were not visited yet. However the algorithm complexity is still bound by $O(|V| + |T|)$.

## 4.3 Generic SPARQL Queries

So far, we only discussed how Lusail evaluates conjunctive SPARQL queries. However, Lusail also supports queries with joins on variable predicates, UNION, FILTER, LIMIT, and OPTIONAL statements, see the listed queries in the Appendix. In a nutshell, Lusail determines where to add clauses, such as FILTER, LIMIT, and OPTIONAL, in the subqueries or during the global join evaluation. For example, FILTER statements with a single variable are pushed with relevant subqueries and thus handled by the endpoints. In case of multi-variable filters, if both variables exist exclusively in the same subquery, then they are handled by the endpoints. Otherwise, Lusail considers the filter clause during the join evaluation phase. Queries with optional statement are handled similarly. If the optional pattern exist within a single subquery and has nothing to do with a global join variable, then the endpoints evaluate it. Otherwise, Lusail takes care of the optional results when finalizing the query answer. Limit and Union statements have naive implementation in Lusail. For queries with a limit clause, SAPE stops the evaluation of the last subquery once the required number of results is reported. Queries with union are split into multiple queries which are evaluated independently and their answers are concatenated after removing the duplicates. Finally, if the query has variable predicate then there are two cases; (*i*) if no predicate join is required or the predicate join exists exclusively within a single subquery, then the triple pattern is pushed to the relevant endpoints for evaluation and no special handling is required. Otherwise, (*ii*) If there is a predicate join across several subqueries, then SAPE handles it during the join evaluation phase.

## 4.4 Result Completeness

**Missing Results.** The optimization introduced by LADE assigns triple patterns to different subqueries based on the concept of locality. Results might be missed in two cases.

**Case 1**: a subquery contains a set of triple patterns where a GJV is considered to be local. This may only happen if the subquery contains triple patterns that access predicates through interlinks, e.g., a subquery that contains ⟨?P, ub:PhDDegreeFrom, ?U⟩ and ⟨?U, ub:address, ?A⟩ will miss the tuple (Kim, Tim, MIT, "XXX") when submitted to EP2. However, such cases cannot happen since LADE puts triple patterns into the same subquery only if the data instances matching them are located in the same endpoint. Lemma 1 states that LADE guarantees that no results are missed.

**Case 2**: a subject or object may be present in more than one endpoint, e.g., EP1 has ⟨a1, p, b⟩, ⟨b, q, c1⟩ and EP2 has ⟨a2, p, b⟩, ⟨b, q, c2⟩. Having the pair of triple patterns, ⟨?x, p, ?y⟩ and ⟨?y, p, ?z⟩, in the same subquery does not miss the local triples matching the query. Lusail first detects $?y$ as a local join variable and then performs the join between the results of the same subquery from different endpoints at the mediator or global level (see Section 5.2).

LEMMA 1. *Any local join variable detected by LADE is a true local join variable.*

PROOF: Let $v$ be the join variable and $TP(v) = \{tp_1, tp_2, ...tp_k\}$ be the set of triple patterns in which $v$ appears. $v$ can appear in $TP(v)$ as subject only, object only, or subject and object.

*Subject only:* In this case, $\forall_{tp_i \in TP(v)}\ tp_i.subj = v$. Let $B_i$ and $B_j$ be the set of bindings of $v$ from triples $tp_i$ and $tp_j$, respectively. LADE decides that $v$ is a local join variable iff: $\forall_{0<l<t}\ \forall_{0<i,j<k, i \neq j}\ B_i(ep_l) - B_j(ep_l) = \phi$ and $B_j(ep_l) - B_i(ep_l) = \phi$ where $k = |TP(v)|$ and $t$ is the number of relevant endpoints. At each relevant endpoint, $B_i - B_j = \phi$ means that each endpoint can fully evaluate $tp_i \bowtie tp_j$ locally. It also means that $\forall_{0<l,r<t,l \neq r} B_i(ep_l) \cap B_j(ep_r) = \phi$, i.e., $v$ is a true local join variable and there is no need to join $tp_i$ and $tp_j$ across endpoints. The same applies for $B_j - B_i$.

*Object only:* In this case, $\forall_{tp_i \in TP(v)}\ tp_i.obj = v$. The same analysis of the subject only case applies.

*Subject/Object:* Let $TPS(v) = \{tps_1, ...tps_s\}$ and $TPO = \{tpo_1, ...tpo_o\}$ be the set of triples in which $v$ appears as subject and object, respectively. $\forall_{tps_i \in TPS(v)}\ tps_i.subj = v$ and $\forall_{tpo_i \in TPO(v)}\ tpo_i.obj = v$. Let $B_i$ and $B_j$ be the set of bindings of $v$ using triple $tps_i$ and $tpo_j$, respectively. LADE decides that $v$ is a local join variable iff: $\forall_{0<l<t}\ \forall_{0<i<s,\ 0<j<o}\ B_i(ep_l) - B_j(ep_l) = \phi$. At each relevant endpoint, $B_i - B_j = \phi$ means that each endpoint can fully evaluate $tps_i \bowtie tpo_j$ locally. It also means that $\forall_{0<l,r<t,l \neq r} B_i(ep_l) \cap B_j(ep_r) = \phi$, i.e., $v$ is a true local join variable and no need to join $tps_i$ and $tpo_j$ across endpoints. □

**Extraneous computations.** In some cases, LADE may detect a join variable as being global while the triple patterns sharing this variable could be solved together locally at the endpoints. For example, the variable ?P in ⟨?S, ub:advisor, ?P⟩ and ⟨?P, ub:teacherOf, ?C⟩. As shown in EP1 (Figure 5), there is an advisor (Ann) who has joined MIT but who is not a teacher of any course yet. ?P will be considered as a GJV based on our checks. However, it is clearly safe to send both triple patterns in the same subquery since there is no need to access data in remote endpoints. Adding more checks for this case would be too expensive since it will require accessing all other relevant endpoints. Such cases may lead to query plans with unnecessary GJVs, i.e., more remote

requests and more join computations at global level rather than at the endpoints. Lemma 2 shows that assuming that a join variable is global, while it is not, does not affect the correctness of the results.

LEMMA 2. *Any local join variable $v$ can be selected as a global join variable without affecting the result correctness.*

PROOF: Let $TP(v) = \{tp_1, ..., tp_k\}$ be the set of triple patterns in which $v$ appeared. If $v$ is a local join variable, each relevant endpoint can evaluate $TP(v)$ as a single subquery. The set of bindings of the local join variable $v$ is simply the union of all bindings from all relevant endpoints, i.e. $B_l(v, TP(v)) = \cup_{0 < i < t} B_l(v, TP(v), ep_i)$ where $t$ is the number of relevant endpoints. Assume now that $v$ is considered a global join variable. In this case, each endpoint will evaluate each triple pattern independently and the results are joined at the global level. Let $B_g(v, tp_j) = \cup_{0 < i < t} B_g(v, tp_j, ep_i)$ be the set of bindings of the global variable $v$ for triple pattern $tp_j$. Then, the global bindings of $v$ is $B_g(v, TP(v)) = B_g(v, tp_1) \bowtie B_g(v, tp_2)... \bowtie B_g(v, tp_k)$. Since $v$ should be a local join variable, then the join between different endpoints is always empty. This means that $B_g(v, TP(v)) = \cup_{0 < i < t} B_g(v, tp_1, ep_i) \bowtie B_g(v, tp_2, ep_i)... \bowtie B_g(v, tp_k, ep_i)$ which is equivalent to evaluating all triples in $TP(v)$ as a single subquery and taking the union across the relevant endpoints. Consequently, $B_g(v, TP(v)) = B_l(v, TP(v))$. □

# 5. SELECTIVITY-AWARE EXECUTION

Selectivity-Aware Planning and parallel Execution (SAPE) algorithm is responsible for choosing: (*i*) a good execution order for the subqueries that would balance between the communication cost and the degree of parallelism and (*ii*) a good join order for the subquery results. An overview of SAPE is shown in Figure 8. SAPE estimates the cardinality of the different subqueries and accordingly delays subqueries expected to return large results. Non-delayed subqueries are evaluated concurrently while the delayed ones are evaluated serially using bound joins. The objective is to maximize the degree of parallelism and to minimize the communication cost in terms of the number of HTTP requests and subquery results. To avoid an increase in the number of HTTP requests due to bound joins, we simply group these bindings into blocks and submit one HTTP request per block.

## 5.1 Subquery Ordering and Cost Model

LADE outputs a set of independent subqueries that can be submitted concurrently for execution at each of its relevant endpoints. The results of these subqueries will then need to be joined at the global level. There are two extreme approaches to execute these subqueries.

The simplest approach is to simultaneously submit the subqueries to the relevant endpoints and wait for their results to start the joining phase. For example, the subqueries of Figure 7 will be executed concurrently and after receiving all their results, a join phase is started. Notice that the subquery ⟨`?U`, address, `?A`⟩ is so generic and executing it independently will retrieve all entities with addresses regardless of whether these entities match `?U` in the remaining subqueries, see Figure 1. Such case is evident in subqueries that touch most of the endpoints or retrieve large amounts of intermediate results.



**Figure 8: Query Evaluation in Lusail**

Such subqueries affect the query evaluation time by overwhelming the network, the endpoints, and Lusail, with irrelevant data. Examples include: (*i*) generic subqueries that are relevant to the majority of the endpoints, e.g., common predicates such as *owl:sameAs*, *rdf:type*, *rdfs:label* and *rdfs:seeAlso*. (*ii*) Simple subqueries that have one triple pattern with two or three variables, e.g., ⟨`?s`, `?p`, `?o`⟩ or ⟨`?s`, owl:sameAs, `?o`⟩, and (*iii*) optional subqueries.

At the other extreme, we submit the most selective subquery first and use the actual bindings of the variables that have been already obtained to submit the next most selective subquery and so on. While limiting the amount of intermediate results to be retrieved from the endpoints, this approach suffers from long communication delays and almost no concurrency except submitting the same subquery to multiple endpoints.

To overcome these shortcomings, we need to balance between the degree of parallelism, i.e., the number of subqueries that can be submitted concurrently, and the communication cost which is dominated by the size of intermediate results. Thus, we detect the subqueries expected to return significantly less results if some of their variables are bound to the already obtained results. The idea is to cluster subqueries based on their estimated cardinality and the number of endpoints they access while taking into account the variability in these values. We introduce the concept of *delayed subqueries*. A delayed subquery is evaluated using the actual bindings of the variables that have been already obtained. We thus follow a two-phase subquery evaluation: (*i*) concurrently submit non-delayed subqueries for evaluation at the endpoints, and (*ii*) use the variable bindings obtained from the first phase to evaluate the delayed subqueries.

We introduce a cost model to determine delayed and non-delayed subqueries. SAPE assumes that subquery cardinalities follow a normal distribution, i.e., most subqueries return results whose sizes are within one standard deviation of the mean. SAPE calculates the average $\mu$ and standard deviation $\sigma$ values for all the cardinalities and for all the numbers of relevant endpoints per subquery. The outliers, e.g., subqueries returning extremely huge results (very low selectivity), or accessing a large number of endpoints, compared to other subqueries, misleadingly increase the standard deviation. This may lead to consider some subqueries, which are better to be delayed, as non-delayed. We therefore apply the Chauvenet's criterion [4] for detecting and rejecting outliers before computing $\mu$ and $\sigma$. Any subquery $sq_i$ with cardinality $C(sq_i) > \mu_C + \sigma_C$ is delayed as shown in Figure 8. We apply the same concept for the number of relevant endpoints per subquery. With this heuristic, only subqueries (including the outliers) whose results are expected to be

significantly bigger than the results of the majority of the subqueries will be delayed.

The cardinality of a subquery is estimated based on the cardinality of its triple patterns. It is collected during the query analysis phase using a simple SELECT COUNT query, one per triple pattern. Whenever a filter clause is available for a subject and/or object, it is pushed with the statistics query to obtain better cardinality estimations. Note that cardinality statistics per predicate are usually collected by RDF engines for their runtime query optimization [8, 20, 15]. Therefore, this count query is lightweight as it asks for a simple triple-by-triple statistics.

We need to estimate the cardinality of the variables in the projection list of each subquery. The cardinality of a variable $v$ in a subquery $sq_i$ is denoted as $C(sq_i, v)$ and represents the number of bindings of $v$. If two triple patterns $TP_i$ and $TP_j$ join on a variable $v$, then the number of bindings of $v$ at endpoint $ep_k$ after the join will be:

$$C(sq_i, v, ep_k) = min(C(TP_i, ep_k), ..., C(TP_j, ep_k))$$

Therefore, we use the minimum cardinality of the predicates, in which $v$ is a common variable, as an upper bound of the cardinality of $v$ per endpoint. Thus, the total cardinality of $v$ in the subquery $sq_i$ is the sum of its cardinalities in all the relevant endpoints $ep$, estimated as:

$$C(sq_i, v) = \sum_{ep \in srcs(sq_i)} C(sq_i, v, ep)$$

The cardinality of a subquery $sq_i$, denoted as $C(sq_i)$, is the maximum cardinality of the subquery projected variables.

## 5.2 Evaluation of Subqueries

Different orders of delayed subquery evaluation can result in different computation and communication costs. Our query planner tries to find an order of subqueries that has the minimal cost. Given a set of non-delayed subqueries, SAPE evaluates them concurrently and builds a hashmap that contains the bindings of each variable. As a result, SAPE knows the exact number of bindings of each subquery variable. Then, we refine the cardinality of the delayed subqueries based on the cardinality of variables they can join with. The first delayed subquery to be evaluated is the one with the lowest cardinality.

Given a set of binding values, SAPE divides them into fixed-size blocks and creates a subquery for each block. The data block size is a SAPE configuration parameter. Once the first subquery is selected, it is evaluated at the corresponding endpoints and its results are used to update the bindings hashmap. SAPE continues to select the next subquery to be evaluated till all subqueries are executed.

Algorithm 3 describes our selectivity-aware evaluation technique for subqueries. The input is a set of independent subqueries with their delay decisions. Each subquery contains its triple patterns, the relevant endpoints (sources), the projection variables, and whether the subquery is optional. The algorithm initializes the request handler which creates a thread per relevant endpoint (line 1). If there is only one subquery; i.e the query is disjoint, the algorithm evaluates the whole query at all relevant endpoints independently (line 3). Then, it aggregates the results obtained from relevant

---

**Algorithm 3:** Subqueries Evaluation

**Input**: Subqueries list ($subQs$), relevant sources ($srcs$)
**Result**: The final query results ($qResult$)

1   $ReqHandler \leftarrow$ initializeRequestHandler ($srcs$);
2   **if** $subQs$.size $()=1$ **then**
3     $ReqHandler$.executeSubQAtRelSrcs ($subQs[0]$);
4     **return** aggregateEndptResults ($ReqHandler$);

5   $foundBindings \leftarrow$ Empty;
6   **foreach** $sq$ in $subQs.nonDelayed$ **do**
7     $ReqHandler$.executeSubQAtRelSrcs ($sq$);

8   $sqsRes \leftarrow$ joinSubqsResults ($ReqHandler.threads$);
9   updateFoundBindings ($subqRes$, $foundBindings$);
10   **while** $subQs.delayed$ $is$ $not$ $empty$ **do**
11     $sq \leftarrow$ getMostSelectiveSubq ($subQs$, $foundBindings$);
12     $boundSubQs \leftarrow$ formulateBoundSubqs ($sq$, $foundBindings$);
13     $sq$.relSrcs $\leftarrow$ refineRelSrs ($sq$.relSrcs, $foundBindings$);
14     $sqRes \leftarrow$ Empty;
15     **foreach** $boundSubq_i$ in $boundSubQs$ **do**
16       $sqRes = sqRes \cup ReqHandler$.executeSubQAtRelSrcs ($boundSubq_i$);
17     updateFoundBindings ($subqRes$, $foundBindings$);
18     $subQs.delayed$.remove ($sq$);
19   **return** joinSubqsResults ($ReqHandler$);

---

endpoints, joins the partial results from different endpoints, if necessary, and returns the final query answer (line 4).

If the query is not disjoint, SAPE iterates over all input subqueries and evaluates each subquery at its relevant endpoints (lines 6-19). In the first phase, non-delayed subqueries are evaluated and their results are collected concurrently (lines 6-7). This step is non-blocking, i.e, each thread is assigned all relevant subqueries at the same time.

Whenever possible, the results of non-delayed subqueries are joined together. This step is necessary to obtain a reduced set of found bindings. In the second phase, SAPE evaluates the delayed subqueries using the found bindings (lines 10-18). SAPE selects the next delayed subquery to be the one with the smallest estimated cardinality (line 11). SAPE formulates a set of modified subqueries from the subquery itself using the found bindings (line 12). It appends a data block to the subquery using the SPARQL VALUES construct, which allows multiple values to be specified in the data block. If the subquery contains triple patterns of the form ⟨?s, ?p, ?o⟩, the source selection process is repeated using the found bindings to reduce the number of relevant endpoints (line 13). Without this refinement, such subqueries are relevant to all endpoints.

We empirically verified that the source selection refinement step on irrelevant endpoints using ASK queries costs significantly less than evaluating the delayed subquery with the found bindings. Finally, the bound subqueries are evaluated and their results are merged (lines 15-16). SAPE updates the set of found bindings using the current subquery results (line 17). After that, the evaluated subquey is removed from the delayed subqueries list (line 18). SAPE continues to evaluate the other subqueries until no more delayed subqueries are left.

**Join Evaluation.** Each endpoint thread maintains a set of relevant subqueries and their corresponding results. This information is encapsulated in the request handler object which is then passed to the threads performing the joins (line 19). Each subquery corresponds to a relation ($R$) for which we know the true cardinality and is partitioned among a set of threads. The join evaluation algorithm has four main steps: ($i$) for each subquery, it collects aggregate statistics

(relation size and number of partitions) from all threads. (*ii*) it then uses a cost-based query optimizer based on the Dynamic Programming (DP) enumeration algorithm [19]. It considers both the inter-operator parallelism where the same operator is being executed concurrently by multiple threads and the selection of the join which leads to the smallest cardinality. Using an in-memory hash join algorithm, joining the subplan at state $S$ with another relation $R$ has two phases; hashing and probing. Assuming that $S$ is the smaller relation, the join cost is estimated as follows:

$$JoinCost(S, R) = \underbrace{\frac{1}{S.threads}|S|}_{hashing} + \underbrace{\frac{1}{R.threads}C(R, v)}_{probing}$$

All threads with the smaller relation build a hash table for their part of $S$. The threads that maintain $R$ will evaluate the join by probing these hash tables with the found bindings of the join variables. (*iii*) Given the devised join order, SAPE joins the different subqueries together to produce the query answer, and (*iv*) finally, SAPE aggregates the joined results from the individual threads and returns the result.

## 6. EXPERIMENTAL STUDY

### 6.1 Evaluation Setup

**Systems Compared:** We evaluate Lusail against: one index-free system, FedX [29], and two index-based systems, SPLENDID [11] and HiBISCuS [26]. [25] has shown that FedX outperformed other systems on the majority of queries and datasets. HiBISCuS [26] is an add-on to improve performance; we use it on top of FedX. SPLENDID showed competitive performance to FedX on several queries in [25] and LargeRDFBench[3]. Similarly to Lusail, both FedX and SPLENDID support multiple-threads.

**Computing Infrastructure:** We used two settings for our experiments: two local clusters, *84-cores* and *480-cores*, and the public cloud. The *84-cores* cluster is a Linux cluster of 21 machines, each with 4 cores and 16GB RAM, connected by 1Gb Ethernet. The *480-cores* cluster is a Linux cluster of 20 machines, each with 24 cores and 148GB RAM, connected by 10Gb Ethernet. We use the *84-cores* cluster in all experiments except those that need 256 endpoints for the LUBM dataset. For the public cloud, we use 18 virtual machines on the Azure cloud to form a real federation.

**Datasets:** We use several real and synthetic datasets. Table 1 shows their statistics. QFed [23] is a federated benchmark of four different real datasets. Although the total number of triples used in QFed is only 1.2 million, there are interlinks between the four datasets, which makes federated query evaluation challenging. LargeRDFBench is a recent federated benchmark of 13 different real datasets with more than 1 billion triples in total. We also used the synthetic LUBM benchmark to generate data for 256 universities, each with around 138K triples. It includes links between the different universities through students and professors.

**Queries:** QFed [23] has different categories of queries. Each query has a label $C$ followed by the number of entities for each class, and a label $P$ followed by the number of predicates linking different datasets. LUBM comes with its benchmark queries. We only used the queries that access multiple endpoints. Queries *Q1*, *Q2*, and *Q3* correspond to

---

[3]https://github.com/AKSW/LargeRDFBench

**Table 1: Datasets used in experiments**

| Benchmark | Endpoint | Triples |
|---|---|---|
| **QFed** | DailyMed | 164,276 |
| | Diseasome | 91,182 |
| | DrugBank | 766,920 |
| | Sider | 193,249 |
| | **Total Triples** | 1,215,627 |
| **LargeRDFBench** | LinkedTCGA-M | 415,030,327 |
| | LinkedTCGA-E | 344,576,146 |
| | LinkedTCGA-A | 35,329,868 |
| | ChEBI | 4,772,706 |
| | DBPedia-Subset | 42,849,609 |
| | DrugBank | 517,023 |
| | Geo Names | 107,950,085 |
| | Jamendo | 1,049,647 |
| | KEGG | 1,090,830 |
| | Linked MDB | 6,147,996 |
| | New York Times | 335,198 |
| | Semantic Web Dog Food | 103,595 |
| | Affymetrix | 44,207,146 |
| | **Total Triples** | 1,003,960,176 |
| **LUBM** | 256 Universities | 35,306,161 |

*Q2*, *Q9*, and *Q13* in the benchmark while *Q4* is a variation of *Q9*, it retrieves extra information from remote universities. LargeRDFBench has three categories: simple *S*, complex *C*, and large *B*. LargeRDFBench subsumes the FedBench benchmark [28]. The complex category contains 10 queries with a high number of triple patterns and advanced SPARQL clauses. The last category has 8 large queries that generate large intermediate data and results. We exclude three queries (*C5*, *B5* and *B6*) as they contain two disjoint subgraphs joined by a filter variable. Neither Lusail nor its competitors support such queries.

**Endpoints:** We use Jena Fuseki 1.1.1 as the SPARQL engine at the endpoints for LUBM and QFed. Since Jena run out of memory while indexing the endpoints of LargeRDFBench, we used a Virtuoso 7.1 instance for each of the 13 endpoints in LargeRDFBench. The standard, unmodified installation of each SPARQL engine is run at the endpoints and used by all federated systems in our experiments.

**Data Preprocessing Cost:** Index-based systems such as SPLENDID and HiBISCuS require a preprocessing phase that generates summaries about the data schemas and collects statistics that are used during query optimization. In real applications, endpoints might not allow collecting these statistics. Moreover, it is a time consuming process dominated by the dataset size. For example, SPLENDID needs 25 and 3,513 seconds to pre-process QFed and LargeRDFBench respectively. In contrast, Lusail and FedX do not require any preprocessing. Hence, index-free methods are preferred in a large scale and dynamic environment, since endpoints can join and leave the federation at no cost.

In the rest of this section, we present the results of our evaluation on a local cluster and on a geo-distributed settings, in Sections 6.2 and 6.3, respectively. We analyze the different costs of Lusail's query processing and its sensitivity to the threshold for delayed queries in Section 6.4. Experiments on multi-query optimization and multi-machine execution are available on Appendices A and B, respectively.

In all subsequent experiments, we report the performance of Lusail with both LADE and SAPE. All systems are allowed to cache the results of source selection. Each query is run three times and we report the average of the last two. We set a time limit of one hour per query before aborting.

### 6.2 Lusail on a Local Cluster
We compare Lusail to FedX, HiBISCuS, and SPLENDID.

Figure 9: Queries with Filter have high selectivity while Big literal queries have bigger intermediate data.



(a) Two Endpoints

(b) Four Endpoints

Figure 10: FedX and HiBISCuS evaluate queries one triple pattern at a time in a bound join. Lusail decomposes these queries based on the location of data instances.



Figure 11: LargeRDFBench: Most of the simple queries do not access large intermediate data, unlike the complex and large queries. X corresponds to time out while missing bars correspond to runtime errors.

They are all deployed on one machine of the *84-cores* cluster. The endpoints are also deployed on the same cluster.

**QFed Dataset:** Figure 9 shows the query performance of Lusail compared to FedX and HiBISCuS. SPLENDID timed out in all QFed queries except *C2P2* which is answered in 56 seconds. Lusail achieves better performance than FedX and HiBISCuS for all queries. Queries with filter, namely *C2P2BF*, *C2P2BOF*, *C2P2F* and *C2P2OF*, have high selectivity, i.e., less intermediate data. Hence, most of these queries are answered within a few seconds. Lusail is up to six times faster than other systems for these queries. Using big literal object (*C2P2B*, *C2P2BO*) increases the volume of communicated data. Hence, FedX and HiBISCuS timed out after one hour in *C2P2BO*, while FedX took significant time to evaluate *C2P2B*, on which HiBISCuS timed out. This is due to the large size of communicated data and the number of remote requests. Lusail successfully answers both these queries in less than 2 seconds.

**LUBM Dataset:** This experiment utilizes up to four university datasets[4] from the LUBM benchmark, each in a different endpoint. Figures 10(a) and 10(b) show the results using two and four endpoints, respectively. The datasets at the endpoints have the same schema. Therefore, FedX and HiBISCuS cannot create exclusive groups. Instead, a subquery is created per triple pattern and is sent to all endpoints. Bound joins are then formulated using all the results retrieved from the different endpoints. This leads to a huge

amount of remote requests. Lusail utilizes the schema as well as the location of data instances accessed by the query to formulate the subqueries. Therefore, Lusail discovered that both *Q1* and *Q2* are disjoint queries and their final results can be formulated by sending the whole query to each endpoint independently.

*Q3* and *Q4* need to join data from different endpoints. *Q3* finds graduate students who received their undergraduate degree from *university0*. This limits the size of intermediate data and the number of endpoints. Hence, FedX and HiBISCuS evaluated this query on two and four different endpoints. Lusail decomposed this query into two subqueries: the first subquery (those who obtained an undergraduate degree from *university0*) is sent to the relevant endpoint. The second subquery contains only ⟨?x, rdf:type, ub:GraduateStudent⟩, which is relevant to all endpoints. Hence, Lusail decided to delay its evaluation and managed to outperform the other systems on four endpoints. Lusail decomposes *Q4* into two subqueries, with the second subquery delayed until the results of the first subquery are ready. Lusail decides the join order of the subquery results based on the cardinality of each result and the number of threads hosting this result (in order to achieve a high degree of parallelism). The figures illustrate that Lusail is up to three orders of magnitude faster than FedX and HiBISCuS for queries *Q1*, *Q2*, and *Q4*. FedX and HiBISCuS ran out of memory for *Q1* on four endpoints. SPLENDID managed to run only *Q3* on four endpoints within 52 seconds,

---

[4]The competitors do not scale to more than four while Lusail (Figures 13(b) and 13(c)) scales to 256 universities.

(a) LargeRDFBench: Complex Queries    (b) LargeRDFBench: Large Queries    (c) LUBM Queries: Two Endpoints

**Figure 12: Geo-distributed federation: endpoints are deployed in 7 different regions of the Azure cloud. Communication cost affects all systems, but Lusail can execute all queries and outperforms other systems.**

which is significantly larger than all other systems.

**LargeRDFBench Dataset:** Figure 11 shows the response times for each system on each query category in LargeRDF-Bench. The performances of Lusail and FedX are comparable for most of the simple queries. The collected statistics by HiBISCuS and SPLENDID do not guarantee better response times in all cases even for simple queries. For example, HiBISCuS is much slower than Lusail for *S13* and *S14*, and SPLENDID has the worst performance in *S6*, *S7*, *S9*, and *S14*. In queries *S13* and *S14*, Lusail is significantly faster than FedX and HiBISCuS since these two queries return relatively large intermediate results. In general, the simple queries do not access huge intermediate data and target datasets of different schema. Hence, Lusail's optimizations do not result in performance improvement especially against index-based systems. However, index-based systems require preprocessing and cannot add endpoints on demand.

The complex and large queries have on average a larger number of triple patterns per query and access a larger amount of intermediate data. Lusail achieves significantly better performance than other systems for most of the complex queries (Figure 11). Both FedX and HiBISCuS could not finish on *C1* and *C9* within an hour. SPLENDID evaluated only 5 out of the 10 complex queries. *C2* is a selective query returning 4 results, which explains why all systems have comparable performance. FedX achieved the best performance for *C4* followed by Lusail, while HiBISCuS could not evaluate the query within one hour. *C4* contains a LIMIT clause of 50 results. Lusail current implementation uses a naive approach for the LIMIT clause. It computes all the final results and returns only the top 50 results. FedX cuts short the query execution once the first 50 results are obtained, hence FedX outperformed Lusail in *C4*. SPLEN-DID achieved the best performance in only *C6* on which other systems have comparable performance.

Lusail is by far the best system for all big queries. FedX and HiBISCuS timed out in two queries and returned no results on another two queries. SPLENDID succeeded on only one query (*B2*) and timed out in the rest.

**Summary.** Lusail is the only system that successfully executes all queries of LargeRDFBench, often showing orders of magnitude better performance. It is never significantly outperformed by other systems. In contrast, other systems time out or throw runtime errors in addition to their performance being highly variable and highly unpredictable.

## 6.3   Lusail in a Geo-Distributed Setting

In this section, we evaluate Lusail by simulating a real scenario on the cloud as well as using real endpoints.

**Using the MS Azure cloud:** We simulate a real setting with Lusail and endpoints are deployed on 7 different regions in the US and Europe of the Azure cloud. We used 18 D4 instances (8 Cores, 28 GB memory), 13 for the LargeRDF-Bench endpoints and four for the LUBM and QFed datasets, interchangeably. Lusail and its competitors are deployed on a D5_V2 instance (16 Cores, 56 GB memory) in Central US, while none of the 18 instances is located in Central US.

The communication cost imposed a clear overhead. For QFed, neither FedX nor HiBISCuS were able evaluate most of the queries. FedX finished only *C2P2BF* in 23 seconds, compared to 1.9 seconds for Lusail, while HiBISCuS finished only *C2P2* in 4,477 seconds, compared to 9.5 seconds for Lusail. Figures 12(a) and 12(b) show the query response times of both complex and large queries on LargeRDFBench. We omit the simple queries since they exhibit the same behavior. The high communication overhead affected the runtime of all systems. For complex queries, FedX timed out on two queries and gave runtime errors in two others. HiBISCuS timed out on three queries but did reasonably well in the rest. SPLENDID was able evaluate only five out of the ten complex queries. Lusail outperformed all other systems in almost all complex queries, in some cases by up to two orders of magnitude (*C1* and *C9*). Large queries show the same behavior. Lusail is the only system that returns results (no time out or runtime errors).

Figure 12(c) shows results on two endpoints of the LUBM dataset. Lusail's query response times increased slightly compared to the local cluster (Figure 10(a)). All queries finished in around 1 second. In contrast, both FedX and Hi-BISCuS require more than 1,000 seconds; an order of magnitude compared to their performance on the local cluster. This shows their sensitivity to the communication overhead since they tend to communicate large volumes of data. With four endpoints, FedX and HiBISCuS were able evaluate only Q3 and ran out of memory or timed out in the rest.

**Real Endpoints:** In this experiment, we use Lusail and FedX to query Bio2RDF[5] endpoints. We extracted three representative queries from a real workload: R1, R2 and R3, from the Bio2RDF query log. R1 accesses three endpoints; DrugBank, HGNC and MGI. R2 joins data from Pharm-GKB and OMIM while R3 integrates data from DrugBank

---

[5]http://bio2rdf.org/

(a) Varying Query Complexity  (b) Varying Data Size (LUBM): Q3  (c) Varying Data Size (LUBM): Q4

**Figure 13: Profiling Lusail by varying query complexity, the number of endpoints, and the data size.**



**Figure 14: Evaluating different threshold values for delayed subquery detection.**

and OMIM. This experiment uses a single machine of our *84-cores* cluster as a mediator. FedX was not able to evaluate these queries as it throws several runtime exceptions. On the other hand, Lusail successfully evaluated queries R1, R2 and R3 in 12, 8 and 35 seconds, respectively. This demonstrates that Lusail is capable of answering queries accessing real, independently deployed, endpoints and provides good performance.

## 6.4   Analyzing Lusail

**Profiling Lusail:** Lusail has three phases: source selection, query analysis using LADE, and query execution using SAPE. In this experiment, we profile these phases while varying the query complexity and data size. We use LargeRDFBench queries with different complexities, simple (*S10*), complex (*C4*), and large (*B1*). Lusail is deployed on a single machine of the *84-cores* cluster. The results are shown in Figure 13(a). Source selection and query analysis require a small amount of time compared to query execution, especially for *C4* and *B1*. As expected, the total response time is dominated by the query execution phase. Lusail's query analysis phase is lightweight, requiring less time than the source selection phase in *S10* and *C4*. *B1* requires performing a union operation between two triple patterns and retrieves its data from the endpoints with the largest data sizes. Hence, the query analysis phase takes slightly more time than the source selection phase. In all cases, query analysis does not add significant overhead.

The cost of query processing in Lusail also depends on the number of endpoints and the sizes of the datasets. Therefore, we profiled Lusail while varying the number of endpoints, which also increases the data size. LUBM allows us to increase both endpoints and data size in a systematic way by adding more universities. We deployed 256 univer-

sity endpoints on the *480-cores* cluster. Lusail is deployed on one machine in the same cluster.

Figures 13(b) and 13(c) show the time required for each phase of *Q3* and *Q4*, respectively. Both queries join data from different endpoints to produce the final result. Lusail's query analysis is lightweight, especially for *Q3* since it has only two triple patterns. For *Q3*, Lusail detects the GJVs using the source selection information, i.e., it does not need to communicate with the endpoints. Source selection time is substantial for these queries and increases slightly as the number of endpoints increases. Query execution time is the dominant factor as the number of endpoints increases. The figure shows the total query response time with and without caching the results of *ASK* and *check* queries. The cache helps, especially for the more complex *Q4* and when the number of endpoints is large.

**Delayed Subqueries:** This experiment evaluates different values for the threshold used to delay subqueries, these are $\mu$, $\mu + \sigma$, and $\mu + 2\sigma$, in addition to delaying only subqueries with outlier estimated cardinalities. We used the Chauvenet criterion [4] for outlier detection. In this experiment, we use our LargeRDFBench deployment in Microsoft Azure. We report the total time for evaluating the queries of each category in LargeRDFBench. Figure 14 shows that for *simple* and *complex* queries, $\mu + 2\sigma$ and *outliers* allowed most subqueries to be evaluated concurrently and only few of them to be delayed. Hence, they missed the opportunity to delay some subqueries that could reduce the communication cost and hence the cost of joining the fetched data. Thus, $\mu + 2\sigma$ and *outliers* achieved significantly worse than $\mu$ and $\mu + \sigma$. For *large* queries, delaying several subqueries limits the parallelism. Thus, $\mu$ achieved significantly worse than others as fewer subqueries were evaluated concurrently while the rest were delayed. As shown, $\mu + \sigma$ consistently performs well in all the three categories and hence we use it in our system.

**Analyzing Lusail Optimizations** This experiment measures the gain obtained by the optimizations introduced by LADE and SAPE compared to FedX. We compare the two systems in terms of memory, network, and response time. FedX and Lusail are each deployed on a single machine of the *84-cores* cluster. We only report results for two queries from each dataset; we observed the same behavior in most of the queries with medium and high complexity.

Figures 15(a), 15(b), and 15(c) show the peak memory usage, communication volume (intermediate data to be shipped), and total response time of each query, respectively. FedX takes a significant amount of time for query execution

**(a) Memory Footprint**   **(b) Communication Volume**   **(c) Response Time**

**Figure 15: Lusail optimizations can move computations on intermediate data to the endpoints, retrieve less irrelevant data, and achieve orders of magnitudes speedup compared to FedX. X corresponds to time out.**

due to its static query decomposition and bound join evaluation. It could not process three queries out of six within the time limit of one hour (*C2P2O*, *C9* and *B1*). In these three queries, FedX did not consume a lot of memory or send a lot of data over the network, but rather it overwhelmed the endpoints with lots of requests and wasted the whole hour waiting for them to finish.

LADE decomposition shifts the computation on the intermediate data from Lusail to the endpoints. Therefore, for the queries that FedX was able to complete, Lusail with LADE alone (without SAPE) consumes significantly less memory and communication compared to FedX. Lusail outperforms FedX by up to three orders of magnitude in terms of query response time. Using the SAPE execution in addition to the LADE decomposition further improved the query response time of Lusail. When enabling SAPE, the delayed queries affect the memory and communication costs, sometimes positively and sometimes negatively. However, communication is performed in parallel using multiple threads. Thus, the net effect is that SAPE with LADE always improve on the query response time compared to LADE alone.

In all subsequent experiments, we report the performance of Lusail with both LADE and SAPE. All systems are allowed to cache the results of source selection. Each query is run three times and we report the average of the last two. We set a time limit of one hour per query before aborting.

## 7. RELATED WORK

Distributed (i.e., parallel) RDF systems [15, 13, 31, 17, 16] deal with data stored in a single endpoint, where the data is replicated and/or partitioned among different servers in the same cluster. The goal of these systems is to speed up query execution for RDF data at one endpoint. In contrast, federated RDF systems have no control over the data; the data is accessible through independent, remote SPARQL endpoints.

Federated SPARQL systems can be classified into index-based and index-free. The source selection in index-based systems, such as ANAPSID [2], SPLENDID [11], DARQ [22], and HiBISCuS [26], is based on collected information and statistics about the data hosted by each endpoint. Therefore, the cost of adding a new endpoint is proportional to the size of the data. Index-free systems, such as Fedx [29] and Lusail, do not assume any prior knowledge of the datasets. Fedx [29] and Lusail utilize SPARQL ASK queries to find the relevant endpoints and cache their results for future usage. Thus, the startup cost and the cost of adding a new endpoint is small.

Federated SPARQL systems usually divides the query into exclusive groups of triple patterns, where each group has a solution at only one endpoint. They do not check whether the data instances matching a group of triples are located in the same endpoint. Thus, they fail to create groups for triple patterns having solutions in different endpoints. Lusail detects whether the data instances are located in disjoint groups, i.e., each endpoint can solve the subquery locally, or distributed groups, i.e., some instances are located in remote endpoints. This allows Lusail to shift most of the computation on intermediate data to the endpoints.

Several efforts, such as Ariadne [3], InfoMaster [7], Garlic [24], and Disco [30], have focused on web-based data integration over heterogeneous information sources [21]. In general, a wrapper is run at each data source to translate between the supported languages and data models. Moreover, systems, such as Piazza [14], coDB [9] and HePToX [5], are peer-to-peer systems that interconnect a network of heterogeneous data sources. Since Lusail works with SPARQL endpoints, it does not need wrappers, and it takes advantage of the capabilities of SPARQL (e.g., ASK). Moreover, while these systems utilize source descriptions (schema), Lusail does not assume any prior knowledge about the datasets. In terms of query decomposition, these systems also aim at dividing a query into exclusive subqueries based on the known schema, where each subquery is submitted to only one data source. In contrast, Lusail's decomposition benefits from the actual location of data matching the query to maximize the local computation and increase parallelism.

## 8. CONCLUSION

Lusail optimizes federated SPARQL query processing through a locality-aware decomposition at compile time followed by selectivity-aware and parallel query execution at run time. The decomposition is based not on the schema but rather on the actual location of data instances satisfying the query triple patterns. This decomposition increases parallelism and minimizes the retrieval of unnecessary data. Selectivity-aware and parallel query execution orders queries at run time by delaying subqueries expected to return large results, and chooses join orders that achieve a high degree of parallelism. Lusail outperforms state-of-the-art systems by orders of magnitude and scales to more than 250 endpoints with data sizes up to billions of triples.

# 9. REFERENCES

[1] I. Abdelaziz, E. Mansour, M. Ouzzani, A. Aboulnaga, and P. Kalnis. Query optimizations over decentralized RDF graphs. In *Proc. Int. Conf. on Data Engineering (ICDE)*, pages 139–142, 2017.

[2] M. Acosta, M.-E. Vidal, T. Lampo, J. Castillo, and E. Ruckhaus. ANAPSID: an adaptive query processing engine for SPARQL endpoints. In *Proc. Int. Semantic Web Conference (ISWC)*, 2011.

[3] J. L. Ambite and C. A. Knoblock. Flexible and scalable query planning in distributed and heterogeneous environments. In *Proc. Int. Conf. on Artificial Intelligence Planning Systems (AIPS)*, 1998.

[4] L. Bol'shev and M. Ubaidullaeva. Chauvenet's test in the classical theory of errors. *Theory of Probability & Its Applications*, 19(4):683–692, 1975.

[5] A. Bonifati, E. Chang, T. Ho, L. V. Lakshmanan, R. Pottinger, and Y. Chung. Schema mapping and query translation in heterogeneous P2P XML databases. *The VLDB Journal*, 19(2), 2010.

[6] D. DeWitt and J. Gray. Parallel database systems: The future of high performance database systems. *Communications of the ACM*, 35(6), 1992.

[7] O. M. Duschka and M. R. Genesereth. Query planning in infomaster. In *Proc. ACM Symposium on Applied Computing (SAC)*, 1997.

[8] O. Erling and I. Mikhailov. RDF support in the Virtuoso DBMS. In *Networked Knowledge-Networked Media*, pages 7–24. Springer, 2009.

[9] E. Franconi, G. Kuper, A. Lopatenko, and I. Zaihrayeu. Queries and updates in the coDB peer to peer database system. In *Proc. Int. Conf. on Very Large Data Bases (VLDB)*, 2004.

[10] M. A. Gallego, J. D. Fernández, M. A. Martínez-Prieto, and P. de la Fuente. An empirical study of real-world SPARQL queries. In *1st International Workshop on Usage Analysis and the Web of Data (USEWOD) at Proc. World Wide Web Conf. (WWW)*, 2011.

[11] O. Görlitz and S. Staab. SPLENDID: SPARQL endpoint federation exploiting VOID descriptions. In *Proc. Workshop on Consuming Linked Data (COLD) at (ISWC)*, 2011.

[12] Y. Guo, Z. Pan, and J. Heflin. LUBM: A benchmark for OWL knowledge base systems. *Web Semantics*, 3(2-3), 2005.

[13] S. Gurajada, S. Seufert, I. Miliaraki, and M. Theobald. TriAD: A Distributed Shared-nothing RDF Engine Based on Asynchronous Message Passing. In *Proc. Int. Conf. on Management of Data (SIGMOD)*, 2014.

[14] A. Y. Halevy, Z. G. Ives, P. Mork, and I. Tatarinov. Piazza: Data management infrastructure for semantic Web applications. In *Proc. Int. Conf. on World Wide Web*, WWW, 2003.

[15] R. Harbi, I. Abdelaziz, P. Kalnis, N. Mamoulis, Y. Ebrahim, and M. Sahli. Accelerating SPARQL queries by exploiting hash-based locality and adaptive partitioning. *The VLDB Journal*, pages 1–26, 2016.

[16] J. Huang, D. Abadi, and K. Ren. Scalable SPARQL Querying of Large RDF Graphs. *The Proceedings of the VLDB Endowment (PVLDB)*, 4(11), 2011.

[17] K. Lee and L. Liu. Scaling Queries over Big RDF Graphs with Semantic Hash Partitioning. *The Proceedings of the VLDB Endowment (PVLDB)*, 6(14), 2013.

[18] E. Mansour, I. Abdelaziz, M. Ouzzani, A. Aboulnaga, and P. Kalnis. A demonstration of Lusail: Querying linked data at scale. In *Proc. Int. Conf. on Management of Data (SIGMOD)*, pages 1603–1606, 2017.

[19] G. Moerkotte and T. Neumann. Analysis of two existing and one new dynamic programming algorithm for the generation of optimal bushy join trees without cross products. In *Proc. Int. Conf. on Very Large Data Bases (VLDB)*, pages 930–941, 2006.

[20] T. Neumann and G. Weikum. The RDF-3X engine for scalable management of RDF data. *The VLDB Journal*, 19(1):91–113, 2010.

[21] M. Ouzzani and A. Bouguettaya. Query processing and optimization on the Web. *Distributed and Parallel Databases*, 15(3), 2004.

[22] B. Quilitz and U. Leser. Querying distributed RDF data sources with SPARQL. In *Proc. European Semantic Web Conference on The Semantic Web: Research and Applications (ESWC)*, 2008.

[23] N. A. Rakhmawati, M. Saleem, S. Lalithsena, and S. Decker. QFed: Query set for federated SPARQL query benchmark. In *Proc. Int. Conf. on Information Integration and Web-based Applications (iiWAS))*, 2014.

[24] M. T. Roth and P. M. Schwarz. Don't scrap it, wrap it! a wrapper architecture for legacy data sources. In *Proc. Int. Conf. on Very Large Data Bases (VLDB)*, volume 97, 1997.

[25] M. Saleem, Y. Khan, A. Hasnain, I. Ermilov, and A.-C. Ngonga Ngomo. A fine-grained evaluation of SPARQL endpoint federation systems. *Semantic Web Journal*, 7(5):493–518, 2015.

[26] M. Saleem and A. N. Ngomo. HiBISCuS: Hypergraph-based source selection for SPARQL endpoint federation. In *Proc. Extended Semantic Web Conference (ESWC)*, 2014.

[27] M. Schmachtenberg, C. Bizer, and H. Paulheim. Adoption of the linked data best practices in different topical domains. In *Proc. Int. Semantic Web Conference (ISWC)*, 2014.

[28] M. Schmidt, O. Görlitz, P. Haase, G. Ladwig, A. Schwarte, and T. Tran. FedBench: A benchmark suite for federated semantic data query processing. In *Proc. Int. Semantic Web Conference (ISWC)*, 2011.

[29] A. Schwarte, P. Haase, K. Hose, R. Schenkel, and M. Schmidt. FedX: Optimization techniques for federated query processing on linked data. In *Proc. Int. Semantic Web Conference (ISWC)*, 2011.

[30] A. Tomasic, L. Raschid, and P. Valduriez. Scaling heterogeneous databases and the design of disco. In *Proc. Int. Conf. on Distributed Computing Systems (ICDCS)*, 1996.

[31] K. Zeng, J. Yang, H. Wang, B. Shao, and Z. Wang. A distributed graph engine for web scale RDF data. *The Proceedings of the VLDB Endowment (PVLDB)*, 6(4), 2013.

# APPENDIX

## A. MULTIPLE-QUERY OPTIMIZATION

We conduct an experiment to evaluate the effectiveness of our MQO add-on feature. We use the LUBM workload in [15]. This workload consists of 10K unique queries which are derived from the 14 LUBM benchmarks by changing their structures and constants. We generate workloads of different sizes by randomly choosing queries from the pool of the 10K unique queries. We generated workloads of 10, 20, 40, 80 and 160 queries. Figure 16 shows the effect of the MQO of Lusail compared to evaluating queries independently. As the workload size increases, the possibility to find similar substructures between queries increases. Therefore, utilizing the shared computation among queries results in performance gains that range from 34% to 46%.



**Figure 16: Effect of MQO for different LUBM workload sizes evaluated on 256 university endpoints.**

## B. MULTI-MACHINE EXECUTION

This experiment evaluates the performance of Lusail as we scale up the amount of data and also the number of cores used by Lusail. We use the LUBM benchmark and vary the number of endpoints (universities) from 2 to 128. We run the experiments on the *480-cores* cluster. The goal of this experiment is to increase the number of cores used by Lusail and demonstrate a good scale up [6], also known as weak scalability. Specifically, we want to demonstrate that Lusail can maintain a fixed query response time as the problem size (number of endpoints) and the number of cores expand simultaneously. This type of scalability is important for federated systems, since it corresponds to being able to handle more endpoints by adding more cores, since each new endpoint comes with its own new data.

For this experiment, we deploy Lusail on 8 machines with 16 threads each. The rest of the machines are used for deploying the endpoints. We double the computing resources (cores) as we double the number of endpoints. Figure 17 shows the results using $Q3$ and $Q4$, which are the non-disjoint queries. The figure shows good scale up for Lusail. The performance for $Q3$ is not as good as $Q4$, since Lusail processes joins by sending the less partitioned relation (inner) to the more partitioned relation (outer) with the aim of increasing parallelism. Both $Q4$ and $Q3$ are decomposed into two subqueries (relations). However, the inner relation in $Q4$ is very small, while the one in $Q3$ is larger. Therefore, sending the inner relation in $Q3$ incurs a higher communication cost.



**Figure 17: Multi-machine Execution: utilizing multiple machines for query evaluation.**

## C. BIO2RDF ENDPOINTS AND QUERIES

The relevant endpoints to the queries R1, R2, and R3 are DrugBank[6], OMIM[7], HGNC[8], MGI[9] and PharmGKB[10].

```
PREFIX dbank: <http://bio2rdf.org/drugbank_
                           vocabulary:>
PREFIX mgi:<http://bio2rdf.org/mgi_vocabulary>
PREFIX hgnc:<http://bio2rdf.org/hgnc_vocabulary>
PREFIX pgkb:<http://bio2rdf.org/pharmgkb_vocabulary>
PREFIX omim:<http://bio2rdf.org/omim_vocabulary>
R1: SELECT ?drug ?hgnc ?model WHERE {
    ?drug   dbank:target ?trgt .
    ?trgt   dbank:x-hgnc ?hgnc .
    ?hgnc   hgnc:x-mgi ?marker .
    ?model mgi:marker ?marker .
    ?model mgi:allele ?all .
    ?all    mgi:allele-attribute ?allele_type .
    FILTER(?drug = <http://bio2rdf.org/drugbank:
                        DB00619>)
}

R2: SELECT ?gene ?ref ?cterm  where{
    <http://bio2rdf.org/pharmgkb:PA446359>
            pgkb:x-snomedct ?cterm .
    ?gene   rdf:type    omim:Gene .
    ?gene   omim:refers-to ?ref .
    ?ref    omim:x-snomed ?cterm .
}

R3: SELECT ?drug ?s ?protein WHERE
{
    ?drug   dbank:gene-name ?geneName.
    ?drug   dbank:x-uniprot ?protein.
    ?drug   dbank:general-function ?genFunction.
    ?drug   dbank:specific-function ?speFunction.
    ?pheno rdf:type omim:Phenotype .
    ?pheno rdfs:label ?o.
    ?pheno omim:clinical-features ?clinicFeature.
    ?pheno omim:article ?article.
    ?pheno omim:x-uniprot ?protein.
} limit 500
```

## D. THE DRUG QUERY BASED ON QFED

---

[6] http://drugbank.bio2rdf.org/sparql

[7] http://omim.bio2rdf.org/sparql

[8] http://hgnc.bio2rdf.org/sparql

[9] http://mgi.bio2rdf.org/sparql

[10] http://pharmgkb.bio2rdf.org/sparql

```
PREFIX foaf: <http://xmlns.com/foaf/0.1/>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
PREFIX drugbank: <http://www4.wiwiss.fu-berlin.de/
        drugbank/resource/drugbank/>
PREFIX dosage: <http://www4.wiwiss.fu-berlin.de/
        drugbank/resource/dosageforms/>
PREFIX sider: <http://www4.wiwiss.fu-berlin.de/
        sider/resource/sider/>
PREFIX dailymed: <http://www4.wiwiss.fu-berlin.de/
        dailymed/resource/dailymed/>
SELECT * WHERE {
 ?disease rdfs:label "Asthma" .
 ?drug drugbank:possibleDiseaseTarget ?disease .
 ?drug  drugbank:dosageForm dosage:tabletOral .
 ?drug  drugbank:brandName ?brand_name .
 ?drug  rdfs:label ?drug_name .
 OPTIONAL {
  ?siderdrug owl:sameAs ?drug .
  ?siderdrug      sider:sideEffect ?sideeffect .
  ?sideeffect rdfs:label "Acidosis" .
  ?moiety rdfs:label ?drug_name .
  ?branded_drug dailymed:activeMoiety ?moiety .
  ?branded_drug  dailymed:contraindication ?contr .
 }
 }
```

## E.  LUBM QUERIES

```
PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-
        syntax-ns#>
PREFIX ub: <http://swat.cse.lehigh.edu/onto/
        univ-bench.owl#>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-
        schema#>
Q1: SELECT ?X ?Y ?Z WHERE{
    ?X rdf:type ub:GraduateStudent .
    ?Y rdf:type ub:University .
    ?Z rdf:type ub:Department .
    ?X ub:memberOf ?Z .
    ?Z ub:subOrganizationOf ?Y .
    ?X ub:undergraduateDegreeFrom ?Y . }

Q2: SELECT ?X ?Y ?Z WHERE{
    ?X rdf:type ub:GraduateStudent .
    ?Y rdf:type ub:AssociateProfessor .
    ?Z rdf:type ub:GraduateCourse .
    ?X ub:advisor ?Y .
    ?Y ub:teacherOf ?Z .
    ?X ub:takesCourse ?Z . }

Q3: SELECT ?X WHERE{
    ?X rdf:type ub:GraduateStudent .
    ?X ub:undergraduateDegreeFrom <www.University0.edu> .
    }

Q4: SELECT ?X ?Y ?Z ?N WHERE{
    ?X rdf:type ub:GraduateStudent .
    ?Y rdf:type ub:AssociateProfessor .
    ?Z rdf:type ub:GraduateCourse .
    ?X ub:advisor ?Y .
    ?Y ub:teacherOf ?Z .
    ?X ub:takesCourse ?Z .
    ?Y ub:undergraduateDegreeFrom ?U .
    ?U ub:name ?N .
}
```

## F.  LARGERDFBENCH QUERIES

We show below samples from each query category of Larg-
eRDFBench[11].

```
PREFIX drgbnk: <http://www4.wiwiss.fu-berlin.de/
```
_____
[11]The full list of queries is available at https://github.com/
AKSW/LargeRDFBench

```
        drugbank/resource/drugbank/>
PREFIX drgcat: <http://www4.wiwiss.fu-berlin.de/
        drugbank/resource/drugbank/drugcategory>
PREFIX dc: <http://purl.org/dc/elements/1.1/>
PREFIX bio: <http://bio2rdf.org/ns/bio2rdf#>
PREFIX biokegg: <http://bio2rdf.org/ns/kegg#>
PREFIX geo: <http://www.w3.org/2003/01/geo/wgs84_pos#>
PREFIX geonames: <http://www.geonames.org/ontology#>
PREFIX owl: <http://www.w3.org/2002/07/owl#>
PREFIX dbpedia: <http://dbpedia.org/ontology/>
PREFIX tcga: <http://tcga.deri.ie/schema/>
PREFIX res: <http://dbpedia.org/resource/>
PREFIX nyt : <http://data.nytimes.com/elements/>
PREFIX chebi: <http://bio2rdf.org/ns/chebi#>
PREFIX mo: <http://purl.org/ontology/mo/>
PREFIX xsd: <http://www.w3.org/2001/XMLSchema#>
PREFIX tcgsch: <http://tcga.deri.ie/schema/>


S3: SELECT ?president ?party ?page WHERE {
    ?president rdf:type dbpedia:President .
    ?president dbpedia:nationality res:United_States .
    ?president dbpedia:party ?party .
    ?x nyt:topicPage ?page .
    ?x owl:sameAs ?president .
}
S11: SELECT ?drugDesc ?cpd ?equation WHERE {
    ?drug drgbnk:drugCategory drgcat:cathartics .
    ?drug drgbnk:keggCompoundId ?cpd .
    ?drug drgbnk:description ?drugDesc .
    ?enzyme biokegg:xSubstrate ?cpd .
    ?enzyme rdf:type biokegg:Enzyme .
    ?reaction biokegg:xEnzyme ?enzyme .
    ?reaction biokegg:equation ?equation .
}
S13: SELECT ?drug ?title WHERE {
   ?drug drgbnk:drugCategory drgcat:micronutrient.
   ?drug drgbnk:casRegistryNumber ?id .
   ?keggDrug rdf:type biokegg:Drug .
   ?keggDrug bio:xRef ?id .
   ?keggDrug dc:title ?title .
}


C2: SELECT ?drug ?keggmass ?chebiIupacName
WHERE
{
    ?drug rdf:type drgbnk:drugs .
    ?drug drgbnk:keggCompoundId ?keggDrug .
    ?keggDrug bio:mass ?keggmass .
    ?drug drgbnk:genericName ?drugBankName .
    ?chebiDrug purl:title ?drugBankName .
    ?chebiDrug chebi:iupacName ?chebiIupacName .
    OPTIONAL {
        ?drug drgbnk:inchiIdentifier ?drugbankInchi .
        ?chebiDrug bio2RDF:inchi ?chebiInchi .
        FILTER (?drugbankInchi = ?chebiInchi)
    }
}

C3: SELECT DISTINCT ?artist ?name ?location
    ?anylocation WHERE {
    ?artist a mo:MusicArtist .
    ?artist foaf:name ?name .
    ?artist foaf:based_near ?location .
    ?location geonames:parentFeature ?locationName .
    ?locationName geonames:name ?anylocation .
    ?nytLocation owl:sameAs ?location.
    ?nytLocation nytimes:topicPage ?news
    OPTIONAL
    {
        ?locationName geonames:name
            'Islamic Republic of Afghanistan' .
```

```
        }
}

C4: SELECT DISTINCT ?countryName ?countryCode
    ?locationMap ?population ?long ?lat ?anthem
    ?fDate ?largestCity ?ethGroup ?motto WHERE {
      ?NYTplace geonames:name ?countryName;
      geonames:countryCode ?countryCode;
      geonames:population  ?population;
      geo:long    ?long;
      geo:lat     ?lat;
      owl:sameAs    ?geonameplace.
      OPTIONAL {
          ?geonameplace dbpedia:capital ?capital;
          dbpedia:anthem ?anthem;
          dbpedia:foundingDate ?fDate;
          dbpedia:largestCity ?largestCity;
          dbpedia:ethnicGroup ?ethGroup;
          dbpedia:motto ?motto.
      }
}
LIMIT 50

B2: SELECT DISTINCT ?patient ?tumorType ?exonValue
WHERE
{
    ?s tcga:bcr_patient_barcode ?patient .
    ?patient tcga:disease_acronym
                <http://tcga.deri.ie/lusc> .
    ?patient tcga:tumor_weight ?weight .
    ?patient tcga:tumor_type ?tumorType .
    ?patient tcga:result ?results .
    ?results tcga:RPKM  ?exonValue .
    FILTER(?weight <= 55)
}

B3: SELECT ?patient ?methylationValue
WHERE
{
    ?s tcga:bcr_patient_barcode ?patient.
    ?patient tcgsch:vital_status "Dead".
    ?patient tcga:bcr_drug_barcode ?drug.
    ?drug tcga:drug_name "Tarceva".
    ?patient tcgsch:age_at_initial_pathologic_
                diagnosis ?age.
    ?patient tcga:result ?results.
    ?results  tcga:beta_value ?methylationValue.
    FILTER(?age <= 51)
}
ORDER BY (?patient)

B7: SELECT DISTINCT ?patient ?p ?o WHERE
{
    ?uri tcga:bcr_patient_barcode ?patient .
    ?patient dbpedia:country ?country.
    ?country dbpedia:populationDensity ?popDensity.
    ?patient tcga:bcr_aliquot_barcode ?aliquot.
    ?aliquot ?p ?o.
    FILTER(?popDensity >= 32)
}
```

## G.  QFED QUERIES

We show below a sample query from QFed benchamrk[12]

```
PREFIX sider: <http://www4.wiwiss.fu-berlin.de/
         sider/resource/sider/>
PREFIX diseasome: <http://www4.wiwiss.fu-berlin.de/
         diseasome/resource/diseasome/>
prefix owl: <http://www.w3.org/2002/07/owl#>
C2P2BOF:
```

_____

[12]The full list of queries is available at https://github.com/
nurainir/QFed

```
SELECT * {
    ?s1 a sider:side_effects .
    ?s1 owl:sameAs ?s2 .
    ?s2 diseasome:associatedGene ?URI .
    OPTIONAL {
      ?s2 diseasome:classDegree ?LITERAL .
    }
    FILTER(?LITERAL >= 3) .
}
```